



Free Questions for DP-700

Shared by Hardin on 06-03-2026

For More Free Questions and Preparation Resources

[Check the Links on Last Page](#)



Question 1

Question Type: MultipleChoice

You have a Fabric workspace that contains a warehouse named Warehouse1.

You have an on-premises Microsoft SQL Server database named Database1 that is accessed by using an on-premises data gateway.

You need to copy data from Database1 to Warehouse1.

Which item should you use?

Options:

- A- an Apache Spark job definition
- B- a data pipeline
- C- a Dataflow Gen1 dataflow
- D- an eventstream

Answer:

B



Explanation:

To copy data from an on-premises Microsoft SQL Server database (Database1) to a warehouse (Warehouse1) in Fabric, a data pipeline is the most appropriate tool. A data pipeline in Fabric is designed to move data between various data sources and destinations, including on-premises databases like SQL Server, and cloud-based storage like Fabric warehouses. The data pipeline can handle the connection through an on-premises data gateway, which is required to access on-premises data. This solution facilitates the orchestration of data movement and transformations if needed.

Question 2

Question Type: MultipleChoice

You need to recommend a solution for handling old files. The solution must meet the technical requirements. What should you include in the recommendation?

Options:

- A- a data pipeline that includes a Copy data activity
- B- a notebook that runs the VACUUM command
- C- a notebook that runs the OPTIMIZE command
- D- a data pipeline that includes a Delete data activity

Answer:

B

Question 3

Question Type: MultipleChoice

You have a Fabric workspace that contains a lakehouse and a notebook named Notebook1. Notebook1 reads data into a DataFrame from a table named Table1 and applies transformation logic. The data from the DataFrame is then written to a new Delta table named Table2 by using a merge operation.

You need to consolidate the underlying Parquet files in Table1.

Which command should you run?



Options:

- A- VACUUM
- B- BROADCAST
- C- OPTIMIZE
- D- CACHE

Answer:

C

Explanation:

To consolidate the underlying Parquet files in Table1 and improve query performance by optimizing the data layout, you should use the OPTIMIZE command in Delta Lake. The OPTIMIZE command coalesces smaller files into larger ones and reorganizes the data for more efficient reads. This is particularly useful when working with large datasets in Delta tables, as it helps reduce the number of files and improves performance for subsequent queries or operations like MERGE.

Question 4

Question Type: MultipleChoice

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have a KQL database that contains two tables named Stream and Reference. Stream contains streaming data in the following format.



Reference contains reference data in the following format.



Both tables contain millions of rows.

You have the following KQL queryset.



You need to reduce how long it takes to run the KQL queryset.

Solution: You change the join type to kind=outer.

Does this meet the goal?

Options:

A- Yes

B- No

Answer:

B

Explanation:

An outer join will include unmatched rows from both tables, increasing the dataset size and processing time. It does not improve query performance.

Question 5

Question Type: MultipleChoice

You have a Fabric workspace that contains a lakehouse named Lakehouse1.

In an external data source, you have data files that are 500GB each. A new file is added every day.

You need to ingest the data into Lakehouse1 without applying any transformations. The solution must meet the following requirements

Trigger the process when a new file is added.

Provide the highest throughput.

Which type of item should you use to ingest the data?

Options:

- A- Event stream
- B- Dataflow Gen2
- C- Streaming dataset
- D- Data pipeline



Answer:

A

Explanation:

To ingest large files (500 GB each) from an external data source into Lakehouse1 with high throughput and to trigger the process when a new file is added, an Eventstream is the best solution.

An Eventstream in Fabric is designed for handling real-time data streams and can efficiently ingest large files as soon as they are added to an external source. It is optimized for high throughput and can be configured to trigger upon detecting new files, allowing for fast and continuous ingestion of data with minimal delay.

Question 6

Question Type: MultipleChoice

You have an Azure key vault named KeyVault1 that contains secrets.

You have a Fabric workspace named Workspace1. Workspace1 contains a notebook named Notebook1 that performs the following tasks:

- * Loads stage data to the target tables in a lakehouse
- * Triggers the refresh of a semantic model

You plan to add functionality to Notebook1 that will use the Fabric API to monitor the semantic model refreshes. You need to retrieve the registered application ID and secret from KeyVault1 to generate the authentication token. Solution: You use the following code segment:

Use notebookutils.credentials.getSecret and specify key vault URL and the name of a linked service.

Does this meet the goal?

Options:



- A- Yes
- B- No

Answer:

B

Question 7

Question Type: MultipleChoice

You have an Azure key vault named KeyVault1 that contains secrets.

You have a Fabric workspace named Workspace1. Workspace1 contains a notebook named Notebook1 that performs the following tasks:

- * Loads stage data to the target tables in a lakehouse
- * Triggers the refresh of a semantic model

You plan to add functionality to Notebook1 that will use the Fabric API to monitor the semantic model refreshes. You need to retrieve the registered application ID and secret from KeyVault1 to

generate the authentication token.

Solution: You use the following code segment:

Use `notebookutils.credentials.getSecret` and specify the key vault URL and key vault secret. Does this meet the goal?

Options:

A- Yes

B- No

Answer:

A

Question 8

Question Type: MultipleChoice

You have a Fabric notebook named Notebook1 that has been executing successfully for the last week. 

During the last run, Notebook1 executed nine jobs.

You need to view the jobs in a timeline chart.

What should you use?

Options:

A- Real-Time hub

B- Monitoring hub

C- the job history from the application run

D- Spark History Server

E- the run series from the details of the application run

Answer:

E

Explanation:

The run series from the details of the application run is the most detailed and relevant feature for visualizing job execution in a timeline format, making it the correct choice for this scenario. It provides an intuitive way to analyze job execution patterns and improve the efficiency of the notebook.



To Get Premium Files for DP-700 Visit

<https://www.p2pexams.com/products/dp-700>

For More Free Questions Visit

<https://www.p2pexams.com/microsoft/pdf/dp-700>

20%
DISCOUNT

